

Artist Classification with Stylometry and Support Vector Machines

Ari Brown, Drew Gleeman, and Mike Micatka

December 19, 2012

1 Introduction

1.1 Problem

The problem with being STEM students is that we have an incredibly poor ability to recognize artists from paintings. However, as with most problems in the modern age, they are easily solved with computers and math. Our specific problem is: given a painting and two possible artists, which artist made the painting? The problem itself is interesting although not important. It approaches the field of forgery detection, as our project is forgery detection at its most basic level: artist classification. We are going about this project by using a machine learning algorithm (Support Vector Machine, SVM) to make an educated estimate on who the artist is. With zero training, the computer would make simply a random guess. To turn this guess into an educated estimate, we train it on data from seven different features of each painting, telling the computer to which artist the features belong to. SVM can only make a decision based on one feature, so we run the SVM once for each feature and then we combine the outcomes of each feature using a weighted voting algorithm, where each feature's result gets a certain number of votes, and the winner is then declared by the program with a "sureness" factor.

1.2 Previous Work

There was a similar project done by undergraduate students at Stanford University in a machine learning class that focused on the machine learning aspect. Blessing and Wen used twelve different features to classify their data, with overlap on five of them with ours. Coincidentally, they chose to use the same machine learning algorithms as we did, which were a Bayesian Analyzer and Support Vector Machines. Our implementation differs from the reference implementation because we use fewer features to base our artist-selection decisions on, and we have different weights for our final decider. Trivially, we also used different data, overlapping with only two of the artists. Our implementation is better than the reference implementation because we make use of a weighted council-like approach to produce a single answer based on

all of our stylometrics. We also use cross validation to get better training and more generalized results.

2 Technical Solution

2.1 Summary

Our program is simple to use. Some time before testing and training are to be done, we run the `produce.features` script. This computationally intensive script generates data for all features for all images and saves it in the `features` folder. Once completed, we run our program. It accepts two artists as input as well as a vector representing the weights of each feature, and returns the correct classification rates for each artist, using cross-validation.

2.2 Data

We chose four artists (Rembrandt, Pollock, Monet, and Picasso) and selected roughly 100 works from each artist. We chose each artist for their fame and peculiarities. Picasso was an interesting artist to choose, simply because his style has ranged from realism as a youth to cubism later in his life. Rembrandt is an artist we had seen mentioned in many papers, so we figured that it would be important to include him in our work. Pollock was chosen as he was an abstract painter with creative uses of colors that focused around a certain theme. Monet was selected as he was an impressionist painter who did many plein-air landscapes.

2.3 Weighted Combination

On a more technical level, the program also takes information on the weights of each feature and the threshold for “sureness”, in which if the resulting sureness fails to meet the sureness, the program will produce a symbol which equates to not knowing. As each individual stylometric feature gives its opinion on which artist produced the painting, the program weights the result as 0 or 1, depending on the determined artist. When all the stylometries have produced a result, the values are then weighted respectively according to the inputted weight matrix. The weighted values are then summed, and the final output is compared with the “sureness” threshold, and a final decision on the artist is produced. This process is completed for each painting. To speed this process, the stylometric data, an invariant, is produced ahead of time and stored in a handwritten file-system database.

2.4 Cross Validation

In the initial implementation of our algorithm, we arbitrarily split our images into two equal-sized groups - one for training and one for testing. The algorithm was then trained on the training data, with its performance being measured on the testing data. However, this naive implementation introduces some weaknesses. Only half of

the data is used for training, so the results of the testing data are biased to reflect this. One common workaround for this problem is known as cross-validation. In *k-fold* cross-validation, the data is divided into k subgroups. Then, each subgroup is used as the testing data exactly once, with all of the rest of the subgroups being used as the training data. The results from the k iterations of the algorithm are then combined (in our case, averaged) in order to generate a more robust statistic for measuring the performance of the algorithm.

2.5 Histogram of Oriented Gradients

The histogram of oriented gradients stylometric, or HoG, is based on the direction of intensity changes in cells across an image. The idea behind it is that changes in intensity mark feature changes, and so a histogram for a small cell is produced based on each pixel. The histograms of each cell are then compiled into one, which is the final stylometric result. Which used a handwritten implementation that was heavily influenced from a different source.

2.6 Edge Detection

Our edge histogram algorithm was an adaptation of the Harris corner detection algorithm that was written for one of the homeworks. The Harris algorithm looks at two eigenvalues for every point, in order to determine if the point is part of a corner, edge, or flat surface. In our version of the algorithm for this, we care only about the larger of the two. For each point, we calculate the eigenvalues and select the largest one. We add that to a list, with one eigenvalue for each point. We then normalize the list by dividing by the largest value in it. All of the new values are then between 0 and 1. The final result of the algorithm is a histogram with four bins, representing the breakdown of the "corner sharpness" for all points in the image. The theory is that an artist who uses well-defined lines will have more points with large eigenvalues, where as an artist with blurry lines will have more low-value points.

2.7 Local Binary Patterns

Local Binary Patterns are patterns that appear in numbers when checking the intensity of individual pixels. It is commonly used in texture identification, so the logic behind using it is that it would identify pixelated artwork such as that which Seurat is famous for (although Seurat's paintings were not used in this project). We used a handwritten implementation.

2.8 Corner Detection

Corner detection simply detects corner in the image. We used a handwritten implementation the detects the corners for a given threshold. The logic behind using this stylometric is that we will be able to tell if an artist prefers to use sharp contrast in two directions in their work. It is an extension of Edge Detection, in that regard.

2.9 Color Histogram

We did not expect much from this stylometric, but wanted to include it because it could shed light on an artist’s overall color intensity. It was calculated by averaging the red, green, and blue values, respectively, for all pixels. The overall intensity of each pixel was also calculated. The motivation behind this was to observe an artist’s overall intensity, as aforementioned, but also to discern any color preferences, however slight.

2.10 SIFT

SIFT or scale-invariant feature transform is an algorithm that is used to describe and locate features in images. The first part of this algorithm involves finding the keypoints which is done by finding the extrema of the Difference of Gaussians (DoG) at different scales.

The SIFT algorithm used here is from the VLFeat library. When run, it finds different features and represents them as a 128-dimensional vector. We combined all of the SIFT features from our training set and clustered them into 10 different groups using k-means clustering. We then generated histograms of frequency for all of the images, both training and testing alike, of each of the feature clusters found in the training images of both artists. We can then use the histograms of the training set to train the SVM and the histograms of the testing images to classify using the trained SVM.

In the VLFeat library, the default `SIFT()` call performs SIFT feature detection for a variety of scale spaces. In the SIFT algorithm, the scale space describes the size of the Gaussian kernel used to smooth the image. Larger scale sizes yield a more smoothed image. In our analysis, we perform SIFT feature detection at a variety of levels in order to more fully analyze the image. We perform the above analysis with a variety of different scale spaces in order to better differentiate between artists. This is useful for art identification because the different styles of artists can be subjectively determined by the patterns that exists at multiple levels, from brushstrokes to the overall theme of a painting. Our algorithm performs SIFT detection with scale spaces ranging from 3 (the default) to 8 (the best-performing version of the algorithm, in our findings). To avoid displaying too much data, we show only the results from scale spaces 3 and 8, in our graphs.

2.11 Blob Detection

Blob detection is a technique in which we calculate the centers of points of interest that have large intensity changes. The radius of a blob measures how far that intensity goes before changing drastically. This means it would be a great stylometric for measuring the presence of overall features that would mark a painting as visually

distracting, such as still lifes and Van Gogh’s famous sunflower painting.

2.12 Fourier Spectral Analysis

Fourier Spectral Analysis is a common technique that was suggested to us by Dr. Dan Rockmore as a replacement for blob detection. Fourier spectral analysis returns Fourier coefficients for the image, where low frequency values represent background and high frequency values represent foreground and texture. The resulting matrix is in the shape of the original image (with the pixels having a change of basis performed on them, so the values do not match up at all), which means it is impossible to train the SVM algorithm on the raw FSA data. As an attempt to solve this, we tried taking the covariance of the data, which also unfortunately returned a matrix that was ultimately based on the size of the original image. We decided not to include Fourier spectral analysis in our final form because we could not figure out a way to faithfully represent the Fourier coefficients, and the overall results were not good enough to be included; they were little more than random guesses.

3 Experiments

In Figure 1, one can see the comparison of each feature’s success in identifying the artist that a painting belonged to. SIFT 8, SIFT 3, and blob detection were the three best individual features, and the overall weighted total also performs just as well as the SIFT 8 does. We ran multiple SIFT versions, but we chose to only display 3 and 8 because 3 was the minimum and 8 was the best. Judging from the fact that corner threshold checking was little better than random guessing, we can determine that these two artists used the same amount of corners and sharp color contrasts. LBP successfully identified many Monet paintings, which would suggest that there are textural differences between the way Monet routinely painted and the way Picasso painted. Picasso may have overlapped Monet’s style, but it was not a routine fact.

Figure 2 shows just how different the two artists are, with almost every stylometry measuring distinct differences and styles. Yet again, SIFT outperformed the rest of the stylometries, and the overall weighting was hardly better than SIFT alone, but the interesting parts here are the results of blob detection, edge histograms, and corner thresholds. Blob detection’s success means that the two artists use very different amounts of visual cues. One artist may have paintings that are more visually distracting than the other’s. The success measured by edge histograms means that the two artists use edges (and, as a subset, corners) very differently in sheer numbers, but the failure of the corner threshold to distinguish Pollock paintings reliably means that corners are used randomly by Pollock and are not a distinguishing point of his style. The success with Monet would suggest that his use of corners is stylistic.

Figure 3 highlights the uniqueness of Rembrandt, as his works are identifiable — in some cases highly so — by every stylometry. Here, the greatest distinguishing




Figure 1: Monet vs Picasso

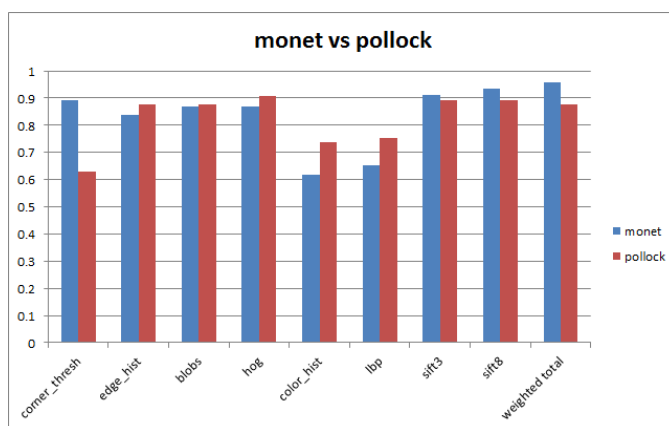


Figure 2: Monet vs Pollock

feature is not the overall weighted metric, but rather the color histogram. The color histogram was suspected to be a poor indicator of artist, but in fact it is proving to be wildly successful for distinguishing Monet and Rembrandt. This would suggest that either each skews wildly to one color, which is unlikely, or that, more likely, Rembrandt's dark tones are setting him apart from Monet.

We can see in Figure 4 that edge histograms were the greatest decider between Picasso and Pollock. It is unusual to note the the overall weighted decider did not pick them as well as any individual feature did (although precedent has certainly been set), but it is more unusual to note that the overall weighted decided was less by so much. LBP failed spectacularly on Picasso, which is unsurprising since Picasso's style




Figure 3: Monet vs Rembrandt

has such range.




Figure 4: Picasso vs Pollock

In Figure 5, we see that SIFT 8 was the greatest decider among the features, and the fact that its superb success rate did not translate into the overall weighted decided means that a majority of the other features disagreed with SIFT at the same time, lending credence to the idea that there are a few specific paintings that did not lend themselves to our analysis. We can see that Rembrandt was more identifiable than Picasso from this graph.

Pollock and Rembrandt, the least identifiable and the most identifiable, in Figure




Figure 5: Picasso vs Rembrandt

6 show us that we can achieve great success rates. Overall, we were able to identify Rembrandt 100% of the time. LBP did the worst out of any stylometry, measuring just under 70% at its worst for Rembrandt and just over 80% for Pollock. Blob detection and edge histograms did phenomenally and outperformed SIFT, which has traditionally been the strongest stylometry. This suggests that Pollock’s abstract paintings are, intuitively, more visually distracting than Rembrandt’s dark scenes.




Figure 6: Pollock vs Rembrandt

4 Conclusions

Overall, we are very pleased with the results we achieved. Our roughly 91% overall identification rate is superior to that of which many individuals are capable. Moreover, we believe that our project lays solid groundwork for anyone looking to tweak or improve our algorithm. Addition of new feature detectors is trivial with the way our program is structured, and changing the weights of each feature is exceedingly simple. Our program also allows for feature data to be computed prior to classification and stored for later use. This allows for maximal preprocessing to be done on a fast machine or for an extended amount of time, and the data to be reused as parameters of the algorithm are tweaked. Below, we will discuss some generalities of the performance of our algorithm. We discuss the across-the-board performance of each feature detector, as well as the identifiability of each individual artist.

Figure 7 is a very telling graph: it tells us what our most reliable measures are, as well as the most unreliable. We can see that, apart from the weighted total, the best-performing algorithms tend to be SIFT, edge histograms, and blobs. Corner threshold and local binary patterns performed the worst, with HoG and color histogram falling somewhere in between. Importantly, we can use these performance measures to gain some sort of idea of how to weight their importance when all of the features are combined. Out of all features, though, the best-performing one is still the weighted total. This is highly desirable (and, in fact, almost necessary), since otherwise selecting just one feature would be ideal. The weights of each feature were initially selected based on the performance of the feature, and then tweaked by hand to maximize the performance.




Figure 7: Overall feature performance

Figure 8 shows how often each artist was correctly identified by the weighted total algorithm. It can be seen that Rembrandt is the most easily identified artist

(with 96% accuracy) while Pollock is the least identifiable (88% accuracy). Monet and Picasso fall somewhere in between (92% and 89%, respectively). It is important to note, however, that these numbers are a function of the weighting chosen in the weighted total algorithm. For instance, Pollock is frequently misidentified by the corner threshold algorithm, so a weight vector favoring corner threshold will yield poor results for Pollock, where it might yield better results for Rembrandt.

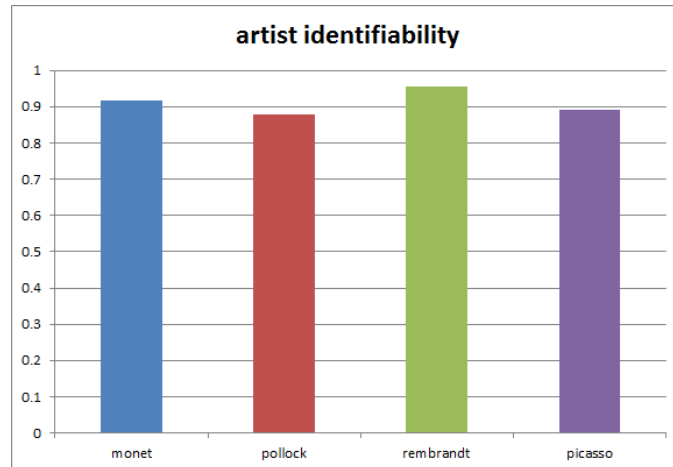


Figure 8: Artist identifiability

References

- [1] Blessing, A. and Wen, K., 2010, *Using Machine Learning for Identification of Art Paintings*.
- [2] Hughes, J. M., Mao, D., Rockmore, D. M., Wang, Y., Wu, Qiang., NOVEMBER 2012, *Empirical Mode Decomposition Analysis for Visual Stylometry*. IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. 34, NO. 11.
- [3] Hughes, J. M., Graham, D. J., Rockmore, D. M., *Stylometrics of artwork: uses and limitations*.
- [4] http://github.com/seyder/art_classifier