

Summer 2012 in Finland

Ari Brown

December 18, 2012

1 Introduction

1.1 Problem

The problem with being STEM students is that we have an incredibly poor ability to recognize artists from paintings. However, as with most problems in the modern age, they are easily solved with computers and math. Our specific problem is: given a painting and two possible artists, which artist made the painting? The problem itself is interesting although not important. It approaches the field of forgery detection, as our project is forgery detection at its most basic level: artist classification.

We are going about this project by using a machine learning algorithm (Support Vector Machine, SVM) to make an educated estimate on who the artist is. With zero training, the computer would make simply a random guess. To turn this guess into an educated estimate, we train it on data from six different features of each painting, telling the computer to which artist the features belong to. SVM can only make a decision based on one feature, so we run the SVM once for each feature and then we combine the outcomes of each feature using a weighted voting algorithm, where each feature's result gets a certain number of votes, and the winner is then declared by the program with a "sureness" factor.

1.2 Previous Work

There was a similar project done by undergraduate students at Stanford University in a machine learning class that focused on the machine learning aspect. Blessing and Wen used twelve different features to

classify their data, with overlap on five of them with ours. Coincidentally, they chose to use the same machine learning algorithms as we did, which were a Bayesian Analyzer and Support Vector Machines.

Our implementation differs from the reference implementation because we use fewer features to base our artist-selection decisions on, and we have different weights for our final decider. Trivially, we also used different data, overlapping with only two of the artists.

Our implementation is better than the reference implementation because we also use Fourier Spectral Analysis which is one of the most basic ways to extract a feature from an image. We also read a paper by Hughes et al. on stylometrics, and another by Hughes et al. on empirical mode decomposition, which inspired our use of Fourier spectral analysis. A more quantitative reason why our implementation is better than the reference implementation is that we achieved better success rates with our system.

2 Technical Solution

2.1 Summary

In its simplest form, our program takes two artists, trains the machine on 40% of the data, and then uses the remaining 60% to test the results while recording the accuracy for estimating each artist. This gives us a measure on both false positives and false negatives, both of which are important.

2.2 Weighted Combination

On a more technical level, the program also takes information on the weights of each feature and the threshold for “sureness”, in which if the resulting sureness fails to meet the sureness, the program will produce a symbol which equates to not knowing. As each individual stylometric feature gives its opinion on which artist produced the painting, the program weights the result as -1 or 1, depending on the determined artist. When all the stylometries have produced a result, the values are then weighted respectively according to the inputted weight matrix. The weighted values are then summed, and the final output is compared with the “sureness” threshold, and a final decision on the artist is produced. This process is completed for each painting. To speed this process, the stylometric data, an invariant, is produced ahead of time and stored in a handwritten file-system database.

2.3 Cross Validation

2.4 Histogram of Oriented Gradients

The histogram of oriented gradients stylometric, or HoG, is based on the direction of intensity changes in cells across an image. The idea behind it is that changes in intensity mark feature changes, and so a histogram for a small cell is produced based on each pixel. The histograms of each cell are then compiled into one, which is the final stylometric result. Which used a handwritten implementation that was heavily influenced from a different source.

2.5 Edge Detection

2.6 Local Binary Patterns

Local Binary Patterns are patterns that appear in numbers when checking the intensity of individual pixels. It is commonly used in texture identification, so the logic behind using it is that it would identify pixelated artwork such as that which Seurat is famous for (although Seurat’s paintings were not used in this project). We used a handwritten implementation.

2.7 Corner Detection

Corner detection simply detects corner in the image. We used a handwritten implementation the detects the corners for a given threshold. The logic behind using this stylometric is that we will be able to tell if an artist prefers to use sharp contrast in two directions in their work. It is an extension of Edge Detection, in that regard.

2.8 Color Histogram

We did not expect much from this stylometric, but wanted to include it because it could shed light on an artist’s overall color intensity. It was calculated by averaging the red, green, and blue values, respectively, for all pixels. The overall intensity of each pixel was also calculated. The motivation behind this was to observe an artist’s overall intensity, as aforementioned, but also to discern any color preferences, however slight.

2.9 SIFT

3 Experiments

4 Conclusions